

Machine Learning for Agricultural Applications

Assignment 2

Prof. Dr. Niels Landwehr
Dr. Julian Adolphs

Summer Term 2020

Release: 07.05.2020
Discussion: 14.05.2020

Task 1 – Linear Regression

[10 points]

Write down the objective function $L(\Theta)$ for linear regression with squared loss and L_2 -Regularization based on the linear model with implicit bias:

$$f_{\mathbf{w}} = \mathbf{w}^T \mathbf{x}, \quad \mathbf{w} \in \mathbb{R}^d, \quad \mathbf{x} \in \mathbb{R}^d.$$

Calculate the gradient of the objective function.

$$L(\Theta) = \frac{1}{n} \sum_{i=1}^n \ell(f_{\Theta}(\mathbf{x}_i), y_i) + \lambda R(\Theta) = \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}_i - y_i)^2 + \lambda \sum_{i=1}^n \mathbf{w}_i^2$$
$$\nabla_{\mathbf{w}} L(\mathbf{w}) = \frac{2}{n} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}_i - y_i) \mathbf{x}_i + 2 \lambda \mathbf{w}$$

Task 2 – Linear Regression, Gradient Descent

[15 points]

- a) Write a python program, that calculates the linear regression with gradient descent for $d = 2$ (i.e. 2 features). Plot the results (the points and the resulting linear function). Generate the points using suitable random generators.
- b) Also calculate the closed form solution and compare with the result from a).

Solution: `ex2_LinReg.py`

Task 3 – Classification

[10 points]

Show that the objective function for classification with cross-entropy loss (without regularization) based on the linear model without bias

$$\mathbf{f}_{\mathbf{W}}(\mathbf{x}_i) = \mathbf{W}\mathbf{x}_i, \quad \mathbf{W} \in \mathbb{R}^{k \times d}, \quad \mathbf{x}_i \in \mathbb{R}^d$$

can be written as

$$L(\mathbf{W}) = -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k \delta_{y_i c_j} \log \left(\frac{\exp(\mathbf{W}_j \mathbf{x}_i)}{\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i)} \right).$$

where \mathbf{W}_j is the j -th row of matrix \mathbf{W} (that is related to class c_j), and $\delta_{y_i c_j}$ is the Kronecker symbol with

$$\delta_{y_i c_j} = \begin{cases} 1, & \text{if } y_i = c_j \\ 0, & \text{if } y_i \neq c_j. \end{cases}$$

Show that the gradient descent for cross-entropy can be calculated iteratively by

$$\mathbf{W}_j^{(k+1)} = \mathbf{W}_j^{(k)} + \alpha \frac{1}{n} \sum_{i=1}^n \left(\delta_{y_i c_j} - \frac{\exp(\mathbf{W}_j \mathbf{x}_i)}{\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i)} \right) \mathbf{x}_i,$$

where $\mathbf{W}_j^{(k)}$ is the k -th iteration step of the j -th row of matrix \mathbf{W} .

Hint: For the differentiation of the loss function, you should take advantage of the logarithm rule $\log\left(\frac{a}{b}\right) = \log a - \log b$. With \log we mean the natural logarithm.

$$\begin{aligned} \ell(f_{\mathbf{W}}(\mathbf{x}_i), y_i) &= -\sum_{j=1}^k \delta_{y_i c_j} \log p(y_i = c_j | \mathbf{x}_i, \mathbf{W}) \\ p(y_i = c_j | \mathbf{x}_i, \mathbf{W}) &= \frac{\exp(f_{\mathbf{W}}(\mathbf{x}_i)_j)}{\sum_{l=1}^k \exp(f_{\mathbf{W}}(\mathbf{x}_i)_l)} \\ \implies \ell(f_{\mathbf{W}}(\mathbf{x}_i), y_i) &= -\sum_{j=1}^k \delta_{y_i c_j} \log \left(\frac{\exp(f_{\mathbf{W}}(\mathbf{x}_i)_j)}{\sum_{l=1}^k \exp(f_{\mathbf{W}}(\mathbf{x}_i)_l)} \right) \\ L(W) &= \frac{1}{n} \sum_{i=1}^n \ell(f_{\mathbf{W}}(\mathbf{x}_i), y_i) \\ \implies L(W) &= -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k \delta_{y_i c_j} \log \left(\frac{\exp(f_{\mathbf{W}}(\mathbf{x}_i)_j)}{\sum_{l=1}^k \exp(f_{\mathbf{W}}(\mathbf{x}_i)_l)} \right) \\ \hookrightarrow L(W) &= -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k \delta_{y_i c_j} \log \left(\frac{\exp(\mathbf{W}_j \mathbf{x}_i)}{\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i)} \right) \quad \text{qed.} \end{aligned}$$

We further convert $L(\mathbf{W})$

$$\begin{aligned} L(\mathbf{W}) &= -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k \delta_{y_i c_j} \left(\log(\exp(\mathbf{W}_j \mathbf{x}_i)) - \log \left(\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i) \right) \right) \\ \implies L(\mathbf{W}) &= -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k \delta_{y_i c_j} \left(\mathbf{W}_j \mathbf{x}_i - \log \left(\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i) \right) \right) \end{aligned}$$

Now we need the partial derivative of $L(\mathbf{W})$:

$$\begin{aligned}
\frac{\partial L(\mathbf{W})}{\partial \mathbf{W}_m} &= -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k \delta_{y_i c_j} \frac{\partial \mathbf{W}_j}{\partial \mathbf{W}_m} \mathbf{x}_i + \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k \delta_{y_i c_j} \frac{\frac{\partial}{\partial \mathbf{W}_m} \sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i)}{\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i)} \\
&= -\frac{1}{n} \sum_{i=1}^n \underbrace{\sum_{j=1}^k \delta_{y_i c_j} \delta_{j m}}_{j=m} \mathbf{x}_i + \frac{1}{n} \sum_{i=1}^n \underbrace{\sum_{j=1}^k \delta_{y_i c_j}}_{=1} \underbrace{\frac{\exp(\mathbf{W}_m \mathbf{x}_i)}{\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i)}}_{\text{because no } j \text{ here}} \mathbf{x}_i \\
&= -\frac{1}{n} \sum_{i=1}^n \delta_{y_i c_m} \mathbf{x}_i + \frac{1}{n} \sum_{i=1}^n \frac{\exp(\mathbf{W}_m \mathbf{x}_i)}{\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i)} \mathbf{x}_i \\
&= -\frac{1}{n} \sum_{i=1}^n \left(\delta_{y_i c_m} - \frac{\exp(\mathbf{W}_m \mathbf{x}_i)}{\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i)} \right) \mathbf{x}_i
\end{aligned}$$

Now we get for the gradient descent steps:

$$\begin{aligned}
\mathbf{W}_j^{(k+1)} &= \mathbf{W}_j^{(k)} - \alpha \frac{\partial L(\mathbf{W})}{\partial \mathbf{W}_j} \\
\hookrightarrow \mathbf{W}_j^{(k+1)} &= \mathbf{W}_j^{(k)} + \alpha \frac{1}{n} \sum_{i=1}^n \left(\delta_{y_i c_m} - \frac{\exp(\mathbf{W}_m \mathbf{x}_i)}{\sum_{l=1}^k \exp(\mathbf{W}_l \mathbf{x}_i)} \right) \mathbf{x}_i \quad \text{qed.}
\end{aligned}$$

Task 4 – Classification, Gradient Descent

[15 points]

Write a python program, that calculates a linear classification with gradient descent for $d = 2$ features and $k = 3$ classes. That can be 3 clusters of points, described by their coordinates (x_1, x_2) and their class labels $y_{\text{label}} \in \{1, 2, 3\}$. Use one-hot-encodings for the class labels $\{1, 2, 3\}$, i.e. class 1: $y = (1, 0, 0)^T$, class 2: $y = (0, 1, 0)^T$ and class 3: $y = (0, 0, 1)^T$. You can use the data given below or create your own data manually or with a random number generator.

```
X = np.array([(-1., 2.5), (-2., 5.), (-1.5, 4.), (-1., 2.3), (-2.5, 6.5), (-1.8, 4.),
              (-1.2, -2.5), (-2.3, -3.), (-1.8, -4.), (-1.9, -2.3), (-2.9, -3.5), (-1.7, -4.),
              (1., -4.5), (0.2, 5.), (0.5, -3.), (1.3, 2.3), (2.5, -1.0), (1.8, 3.)])
```

```
y = np.array([(1, 0, 0), (1, 0, 0), (1, 0, 0), (1, 0, 0), (1, 0, 0), (1, 0, 0),
              (0, 1, 0), (0, 1, 0), (0, 1, 0), (0, 1, 0), (0, 1, 0), (0, 1, 0),
              (0, 0, 1), (0, 0, 1), (0, 0, 1), (0, 0, 1), (0, 0, 1), (0, 0, 1)])
```

Hints: It could be useful to define a function `sigma(x)`. If possible avoid loops, prefer vectorization!

- Plot the progress of the gradient descent by plotting the loss function.
 - Try to visualize the classification boundaries.
 - Use the matrix \mathbf{W} for a class prediction for 5 test points (new points, that are not part of the training set above) and interpret the result.
-

Solution: `ex2_SoftmaxReg.py`

and: `ex2_sklearnClass.py`
